

# Homework 5 Solutions

Due: Thursday 3/26/20 by 8:30am

Rubric:

- Maximum of 2 points each for 1., 2., and 5., determined as follows:
  - 0 points for no solutions whatsoever or incomplete solutions;
  - 1 point for solutions provided for each part, but at least one incorrect solution;
  - 2 points for correct solutions to each part;
- Maximum of 3 points for 3., 4., 6., and 7., determined as follows:
  - 0 points for no solutions whatsoever or R output only;
  - 1 point for an honest effort but very few correct answers or R output only plus a figure;
  - 2 points for mostly correct answers but at least one substantial issue;
  - 3 points for nearly/exactly correct.

1. Problem 2.10 from the .pdf version of the textbook.

- (a) A **prediction** interval is appropriate, because we are interested in the value of a new observation.
- (b) A **confidence** interval is appropriate, because we are interested in the average value of new observations with the same disposable income of \$23,500.
- (c) A **prediction** interval is appropriate, because we are interested in the value of a new observation.

2. Problem 2.11 from the .pdf version of the textbook.

There is a difference - we generally have more uncertainty about one or a finite number of new observations at  $X = X_h$  than about the average across infinitely many observations at  $X = X_h$ , which would be the mean response at  $X = X_h$ .

3. Problem 2.28 from the .pdf version of the textbook, parts (a)-(b).

(a)

```
link <- url("http://maryclare.github.io/stat525/content/homework/muscle.RData")
load(link)
close(link)
Y <- data$Y
X <- data$X

n <- length(Y)

linmod <- lm(Y~X)
b0 <- linmod$coefficients[1]
b1 <- linmod$coefficients[2]
X.h <- 60
Y.hat.h <- b0 + b1*X.h
s.sq <- summary(linmod)$sigma^2

# Construct interval by hand
s.sq.Y.hat.h <- s.sq*(1/n + (X.h - mean(X))^2/(sum((X - mean(X))^2)))
int <- c(Y.hat.h + qt(0.025, n - 2)*sqrt(s.sq.Y.hat.h),
        Y.hat.h + qt(0.975, n - 2)*sqrt(s.sq.Y.hat.h))
# Note: This gives the same result as
```

```
int <- predict(linmod, newdata = data.frame("X" = 60),
              level = 0.95, interval = "confidence")[1, 2:3]
```

We obtain a 95 percent confidence interval of (82.8347139, 87.0589529).

(b)

```
# Construct interval by hand
s.sq.Y.hat.h <- s.sq*(1/n + (X.h - mean(X))^2/(sum((X - mean(X))^2)))
int <- c(Y.hat.h + qt(0.025, n - 2)*sqrt(s.sq + s.sq.Y.hat.h),
        Y.hat.h + qt(0.975, n - 2)*sqrt(s.sq + s.sq.Y.hat.h))
# Note: This gives the same result as
int <- predict(linmod, newdata = data.frame("X" = 60),
              level = 0.95, interval = "prediction")[1, 2:3]
```

We obtain a 95 percent prediction interval of (68.450669, 101.4429978).

4. Problem 2.31 from the .pdf version of the textbook, parts (a)-(b).

(a)

```
link <- url("http://maryclare.github.io/stat525/content/homework/crime.RData")
load(link)
close(link)
Y <- data$Y
X <- data$X

n <- length(Y)

linmod <- lm(Y~X)

ssto <- sum((Y - mean(Y))^2)
sse <- sum(linmod$residuals^2)
ssr <- sum((linmod$fitted.values - mean(Y))^2)

msto <- ssto/(n - 1)
mse <- sse/(n - 2)
msr <- ssr/1
```

Source of Variation	SS	df	MS
Regression	$9.3462942 \times 10^7$	1	$9.3462942 \times 10^7$
Error	$4.5527317 \times 10^8$	82	$5.5521118 \times 10^6$
Total	$5.4873611 \times 10^8$	83	$6.6112784 \times 10^6$

(b)

```
F.star <- msr/mse
F.quantile.01 <- qf(0.99, 1, n - 2)
p.val <- pf(F.star, 1, n - 2, lower.tail = FALSE)
```

To decide whether or not there is a linear association between amount of muscle mass and age, we could conduct an  $F$  test of the null hypothesis  $H_0 : \beta_1 = 0$  against the alternative  $H_a : \beta_1 \neq 0$ . The decision rule for a level- $\alpha = 0.01$  test based on the test statistic  $F^* = MSR/MSE$  would be:

- If  $F^* \leq F(0.99; 1; 82)$ , conclude  $H_0$
- If  $F^* > F(0.99; 1; 82)$ , conclude  $H_a$ .

Because  $F^* = 16.8337645$  and  $F(0.99; 1; 82) = 6.9544199$ , we would reject  $H_0$  and conclude  $H_a$ .

The  $p$ -value of the test is  $P(F > F^*) = 9.5713958 \times 10^{-5}$ , the probability that a  $F$  random variable with 1 and 82 degrees of freedom is greater than  $F^*$  in absolute value. This is identical to the  $p$ -value of the  $t$  test we performed in last week's homework! The decision rules are also equivalent. The decision rule for the  $t$  test was:

- If  $t_{82}(0.005) \leq t^* \leq t_{82}(0.995)$ , conclude  $H_0$
- If  $t^* < t_{82}(0.005)$  or  $t^* > t_{82}(0.995)$ , conclude  $H_a$ .

This can be rewritten as:


- If  $(t^*)^2 \leq t_{82}(0.995)^2$ , conclude  $H_0$
- If  $(t^*)^2 > t_{82}(0.995)^2$ , conclude  $H_a$ .

From last time, we had  $t^* = -4.1028971$ , and  $t_{82}(0.995) = -2.6371234$ . We can see that  $(t^*)^2 = 16.8337645 = F^*$ , and  $t_{82}(0.995)^2 = 6.9544199 = F^*$ , i.e. the test statistics and decision rules are numerically equivalent!

5. Problem 3.2 from the .pdf version of the textbook. The idea is to make an example plot of the residuals  $e_1, \dots, e_n$  against the predictors  $X_1, \dots, X_n$ , that depicts what you might expect to see under the circumstances described by each part of the problem.


(a)

```
X <- 1:10
e <- c(0.30, -0.67, -0.18, 0.35, 0.19,
      0.33, -0.03, -0.15, -0.07, -0.07)
plot(X, e, ylab = "Residuals", pch = 16)
```



(b)

```
X <- 1:10
e <- c(12.34, 4.52, -1.58, -6.09, -8.77,
      -9.85, -7.57, 0.49, 5.40, 11.11)
plot(X, e, ylab = "Residuals", pch = 16)
```



6. Problem 3.7 from the .pdf version of the textbook, parts (a)-(d). You only need to do the first part of (d), i.e. construct a normal probability plot of the residuals. Note that this will require a bit of material that we will cover Tuesday 3/10.

```
link <- url("http://maryclare.github.io/stat525/content/homework/muscle.RData")
load(link)
close(link)
Y <- data$Y
X <- data$X

n <- length(Y)
```

We can use `stem(X, scale = 3)` to create a stem-and-leaf plot.

Stem	Leaf
40	000
42	0000
44	00
46	0000
48	00
50	0
52	000
54	000
56	0000
58	000
60	0000
62	000
64	0000
66	0
68	000
70	000
72	000
74	0


Stem	Leaf
76	0000
78	00000

This plot appears to be consistent with the random selection of women from each 10-year age group, because there are no clear patterns of some ages being better represented in the data than others.

(b)

A dot plot is given below. If you used a bar chart like we did in class, that is ok and will be given full credit, however I think this is clearer.

```
linmod <- lm(Y~X)
e <- linmod$residuals
tab <- table(e)
plot(e, e, xlim = range(e), ylim = c(0, 2),
     xlab = "Residuals",
     ylab = "Number of Units", type = "n")
points(as.numeric(names(tab)),
       tab, pch = 16)
```





The residuals seem evenly distributed about 0, but it is difficult to discern much from a dot plot.

(c)

```
Y.hat <- linmod$fitted.values

par(mfrow = c(1, 2))
plot(X, e, ylab = "Residuals", pch = 16)
plot(Y.hat, e, xlab = expression(hat(Y)),
     ylab = "Residuals", pch = 16)
```




There are no striking relationships between the residuals and the covariate  $X$  or the fitted values  $\hat{Y}$ . There may be slight evidence of increasing variance as  $\hat{Y}$  increases, and of a slight quadratic trend in  $X$ , but these seem minor.

(d)

```
MSE <- sum(e^2)/(n - 2)
Ee <- sqrt(MSE)*qnorm((1:n - 0.375)/(n + 0.25),
                    mean = 0, sd = 1)

plot(Ee, sort(e),
     pch = 16,
     xlab = "Expected",
     ylab = "Residual")
abline(a = 0, b = 1, lty = 3)
```



With the exception of the most extreme residuals, the observed residuals resemble what we would expect to observe if the errors were normally distributed.

7. Problem 3.8 from the .pdf version of the textbook, parts (a)-(d). You only need to do the first part of (d), i.e. construct a normal probability plot of the residuals. Note that this will require a bit of material that we will cover Tuesday 3/10.

(a)

```
link <- url("http://maryclare.github.io/stat525/content/homework/crime.RData")
load(link)
close(link)
Y <- data$Y
X <- data$X

n <- length(Y)
```

We can use `stem(X, scale = 3)` to create a stem-and-leaf plot.


Stem	Leaf
60	0
62	
64	0000
66	00
68	0
70	00
72	00
74	00000000
76	000000000000
78	00000000000000
80	000000000000
82	000000000000000
84	00000000
86	0

Stem	Leaf
88	0000
90	00

The plot indicates that most of the counties represented in the data have high school diploma per capita rates of around 75-80 percent. In every county, at least 60 percent but no more than 90 percent of the residents have a high school diploma.

(b)

```
linmod <- lm(Y~X)
e <- linmod$residuals
boxplot(e,
        horizontal = TRUE,
        xlab = "Residuals")
```




The residuals appear to be slightly skewed to the right, with several very large residuals and generally more positive residuals than negative residuals.

(c)

```
Y.hat <- linmod$fitted.values
plot(Y.hat, e, xlab = expression(hat(Y)),
     ylab = "Residuals")
```






The residuals tend to be larger and have greater variance for nearly average values of  $\hat{Y}$ , which suggests the presence of some slight nonconstant variance or nonlinearity. However overall, there are no striking trends.

(d)

```
MSE <- sum(e^2)/(n - 2)
Ee <- sqrt(MSE)*qnorm((1:n - 0.375)/(n + 0.25),
                    mean = 0, sd = 1)
plot(Ee, sort(e),
     pch = 16,
     xlab = "Expected",
     ylab = "Residual")
abline(a = 0, b = 1, lty = 3)
```



Again, with the exception of the most extreme residuals, the residuals are consistent with what we would expect if the errors were indeed normally distributed.